


Computational Biology

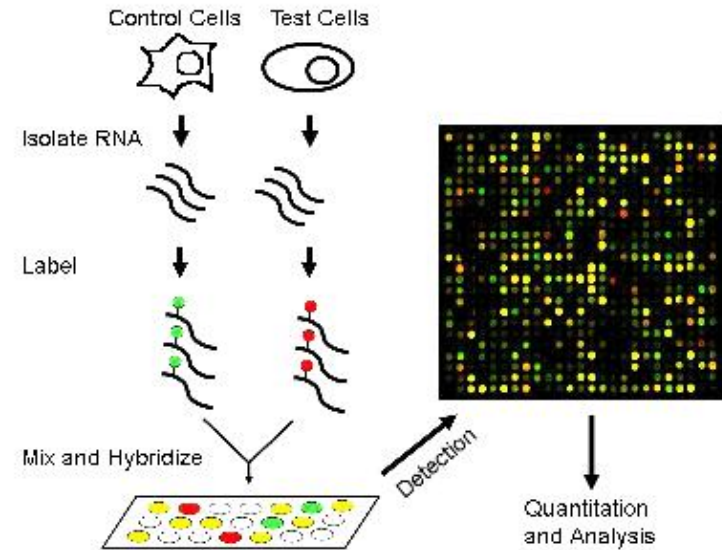
Lecture 2: *Protein Expression*



Microarrays, proteomics, protein interaction

Outline

1. **Protein Expression**
2. Correlations and Causality



Protein Interactions

- Cells are complex organisms where proteins work together to solve problems
- The concentration of proteins will vary depending on what the cell is doing
- The concentration of one protein will affect (excite or inhibit) the expression of other proteins
- This lecture is about try to find these interactions
- We will later look at the interactions in detail

Protein Interactions

- Cells are complex organisms where proteins work together to solve problems
- The concentration of proteins will vary depending on what the cell is doing
- The concentration of one protein will affect (excite or inhibit) the expression of other proteins
- This lecture is about try to find these interactions
- We will later look at the interactions in detail

Protein Interactions

- Cells are complex organisms where proteins work together to solve problems
- The concentration of proteins will vary depending on what the cell is doing
- The concentration of one protein will affect (excite or inhibit) the expression of other proteins
- This lecture is about try to find these interactions
- We will later look at the interactions in detail

Protein Interactions

- Cells are complex organisms where proteins work together to solve problems
- The concentration of proteins will vary depending on what the cell is doing
- The concentration of one protein will affect (excite or inhibit) the expression of other proteins
- This lecture is about try to find these interactions
- We will later look at the interactions in detail

Protein Interactions

- Cells are complex organisms where proteins work together to solve problems
- The concentration of proteins will vary depending on what the cell is doing
- The concentration of one protein will affect (excite or inhibit) the expression of other proteins
- This lecture is about try to find these interactions
- We will later look at the interactions in detail

Proteomics

- The direct measure of protein concentration is known as **proteomics**
- It is a big subject in itself
- Chemically there is no easy way to read off a protein sequence
- We can measure the mass and charge of a molecule accurately using a mass spectrometer
- This is used to directly measure protein concentration

Proteomics

- The direct measure of protein concentration is known as **proteomics**
- It is a big subject in itself
- Chemically there is no easy way to read off a protein sequence
- We can measure the mass and charge of a molecule accurately using a mass spectrometer
- This is used to directly measure protein concentration

Proteomics

- The direct measure of protein concentration is known as **proteomics**
- It is a big subject in itself
- Chemically there is no easy way to read off a protein sequence
- We can measure the mass and charge of a molecule accurately using a mass spectrometer
- This is used to directly measure protein concentration

Proteomics

- The direct measure of protein concentration is known as **proteomics**
- It is a big subject in itself
- Chemically there is no easy way to read off a protein sequence
- We can measure the mass and charge of a molecule accurately using a mass spectrometer
- This is used to directly measure protein concentration

Proteomics

- The direct measure of protein concentration is known as **proteomics**
- It is a big subject in itself
- Chemically there is no easy way to read off a protein sequence
- We can measure the mass and charge of a molecule accurately using a mass spectrometer
- This is used to directly measure protein concentration

Measuring mRNA Levels

- An indirect method to measure protein concentration is to measure the concentration of mRNA
- Although not a perfect proxy it does provide a reasonable approximation to the protein concentration
- It is much easier and cheaper than direct measurement of protein concentration
- Two methods
 - ★ Microarrays (used for the last 20 years)
 - ★ RNA sequencing (beginning to take over)

Measuring mRNA Levels

- An indirect method to measure protein concentration is to measure the concentration of mRNA
- Although not a perfect proxy it does provide a reasonable approximation to the protein concentration
- It is much easier and cheaper than direct measurement of protein concentration
- Two methods
 - ★ Microarrays (used for the last 20 years)
 - ★ RNA sequencing (beginning to take over)

Measuring mRNA Levels

- An indirect method to measure protein concentration is to measure the concentration of mRNA
- Although not a perfect proxy it does provide a reasonable approximation to the protein concentration
- It is much easier and cheaper than direct measurement of protein concentration
- Two methods
 - ★ Microarrays (used for the last 20 years)
 - ★ RNA sequencing (beginning to take over)


Measuring mRNA Levels

- An indirect method to measure protein concentration is to measure the concentration of mRNA
- Although not a perfect proxy it does provide a reasonable approximation to the protein concentration
- It is much easier and cheaper than direct measurement of protein concentration
- Two methods
 - ★ Microarrays (used for the last 20 years)
 - ★ RNA sequencing (beginning to take over)

Measuring mRNA Levels



- An indirect method to measure protein concentration is to measure the concentration of mRNA
- Although not a perfect proxy it does provide a reasonable approximation to the protein concentration
- It is much easier and cheaper than direct measurement of protein concentration
- Two methods
 - ★ Microarrays (used for the last 20 years)
 - ★ RNA sequencing (beginning to take over)

DNA Microarrays



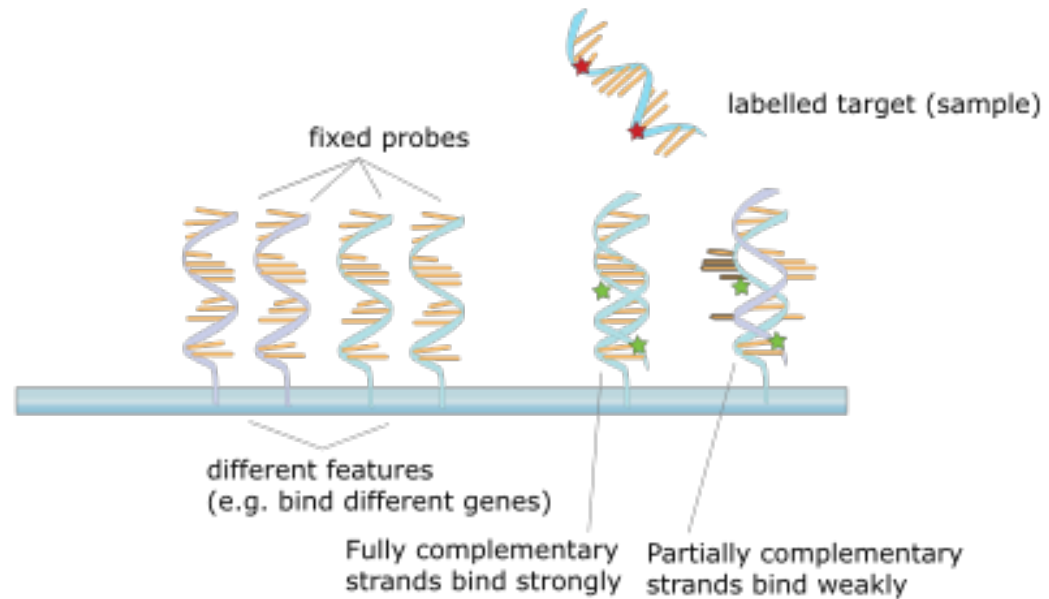
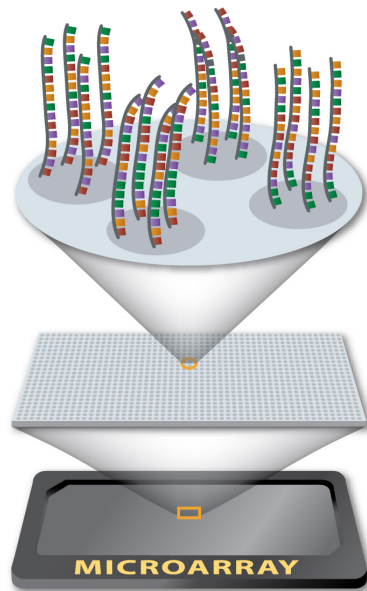
DNA Microarrays

- RNA turned into cDNA
- Micro-wells contain complementary strands of DNA



DNA Microarrays

- RNA turned into cDNA
- Micro-wells contain complementary strands of DNA



DNA Microarrays

- Single stranded mRNA is synthesised to double stranded cDNA
- The cDNA will bind to a complementary strand attached to a microarray
- The microarray needs to be built with the correct DNA strands for the organism being used
- Typically a microarray will have 10 000–100 000 different DNA sequences
- Microarray technology is relatively cheap ($\approx$ £100–£1000)

DNA Microarrays

- Single stranded mRNA is synthesised to double stranded cDNA
- The cDNA will bind to a complementary strand attached to a microarray
- The microarray needs to be built with the correct DNA strands for the organism being used
- Typically a microarray will have 10 000–100 000 different DNA sequences
- Microarray technology is relatively cheap ($\approx$ £100–£1000)

DNA Microarrays

- Single stranded mRNA is synthesised to double stranded cDNA
- The cDNA will bind to a complementary strand attached to a microarray
- The microarray needs to be built with the correct DNA strands for the organism being used
- Typically a microarray will have 10 000–100 000 different DNA sequences
- Microarray technology is relatively cheap ($\approx$ £100–£1000)

DNA Microarrays

- Single stranded mRNA is synthesised to double stranded cDNA
- The cDNA will bind to a complementary strand attached to a microarray
- The microarray needs to be built with the correct DNA strands for the organism being used
- Typically a microarray will have 10 000–100 000 different DNA sequences
- Microarray technology is relatively cheap ($\approx$ £100–£1000)

DNA Microarrays

- Single stranded mRNA is synthesised to double stranded cDNA
- The cDNA will bind to a complementary strand attached to a microarray
- The microarray needs to be built with the correct DNA strands for the organism being used
- Typically a microarray will have 10 000–100 000 different DNA sequences
- Microarray technology is relatively cheap ($\approx$ £100–£1000)

Using DNA Microarray

- All tests are carried out relative to a control cell
- Results show over-expression (red) or under-expression (green) relative to a group of control cells
- Study cell that is doing something different to the control
- Always use a population of cells so we are only looking at population averages

Using DNA Microarray

- All tests are carried out relative to a control cell
- Results show over-expression (red) or under-expression (green) relative to a group of control cells
- Study cell that is doing something different to the control
- Always use a population of cells so we are only looking at population averages

Using DNA Microarray

- All tests are carried out relative to a control cell
- Results show over-expression (red) or under-expression (green) relative to a group of control cells
- Study cell that is doing something different to the control
- Always use a population of cells so we are only looking at population averages

Using DNA Microarray

- All tests are carried out relative to a control cell
- Results show over-expression (red) or under-expression (green) relative to a group of control cells
- Study cell that is doing something different to the control
- Always use a population of cells so we are only looking at population averages

Time Series

- To get an understanding of the correlation between proteins we usually study the evolution over time of the cells
- E.g. add sugar, starve the cell, increase the temperature
- Note, that we are looking at a different population of cells since we kill the cells to extract RNA
- Cells will typically be dividing through the experiment
- Measurements are noisy by their nature

Time Series

- To get an understanding of the correlation between proteins we usually study the evolution over time of the cells
- E.g. add sugar, starve the cell, increase the temperature
- Note, that we are looking at a different population of cells since we kill the cells to extract RNA
- Cells will typically be dividing through the experiment
- Measurements are noisy by their nature

Time Series

- To get an understanding of the correlation between proteins we usually study the evolution over time of the cells
- E.g. add sugar, starve the cell, increase the temperature
- Note, that we are looking at a different population of cells since we kill the cells to extract RNA
- Cells will typically be dividing through the experiment
- Measurements are noisy by their nature

Time Series

- To get an understanding of the correlation between proteins we usually study the evolution over time of the cells
- E.g. add sugar, starve the cell, increase the temperature
- Note, that we are looking at a different population of cells since we kill the cells to extract RNA
- Cells will typically be dividing through the experiment
- Measurements are noisy by their nature

Time Series

- To get an understanding of the correlation between proteins we usually study the evolution over time of the cells
- E.g. add sugar, starve the cell, increase the temperature
- Note, that we are looking at a different population of cells since we kill the cells to extract RNA
- Cells will typically be dividing through the experiment
- Measurements are noisy by their nature

RNA-SEQ

- Recently, sequencing technology has become so cheap that it is feasible to sequence all the mRNAs *de novo*
- This prevents the need for having to tailor the microarray to a specific organism (can move away from model organisms)
- Microarrays can saturate and can provide unreliable counts for poorly expressed proteins due to background noise
- Will find rarely expressed RNA that might not be in microarray

RNA-SEQ

- Recently, sequencing technology has become so cheap that it is feasible to sequence all the mRNAs *de novo*
- This prevents the need for having to tailor the microarray to a specific organism (can move away from model organisms)
- Microarrays can saturate and can provide unreliable counts for poorly expressed proteins due to background noise
- Will find rarely expressed RNA that might not be in microarray

RNA-SEQ

- Recently, sequencing technology has become so cheap that it is feasible to sequence all the mRNAs *de novo*
- This prevents the need for having to tailor the microarray to a specific organism (can move away from model organisms)
- Microarrays can saturate and can provide unreliable counts for poorly expressed proteins due to background noise
- Will find rarely expressed RNA that might not be in microarray

RNA-SEQ

- Recently, sequencing technology has become so cheap that it is feasible to sequence all the mRNAs *de novo*
- This prevents the need for having to tailor the microarray to a specific organism (can move away from model organisms)
- Microarrays can saturate and can provide unreliable counts for poorly expressed proteins due to background noise
- Will find rarely expressed RNA that might not be in microarray

Single Cell RNA-SEQ

- With third generation sequencing technology it is even possible to measure the RNA concentrations in a single cell
- This is expensive, but allows researchers to identify different cells in a population
- Can be very important, for examples, in studying tumours
- Nevertheless, due to amplification processes, etc. there can be biases

Single Cell RNA-SEQ

- With third generation sequencing technology it is even possible to measure the RNA concentrations in a single cell
- This is expensive, but allows researchers to identify different cells in a population
- Can be very important, for examples, in studying tumours
- Nevertheless, due to amplification processes, etc. there can be biases


Single Cell RNA-SEQ

- With third generation sequencing technology it is even possible to measure the RNA concentrations in a single cell
- This is expensive, but allows researchers to identify different cells in a population
- Can be very important, for examples, in studying tumours
- Nevertheless, due to amplification processes, etc. there can be biases

Single Cell RNA-SEQ


- With third generation sequencing technology it is even possible to measure the RNA concentrations in a single cell
- This is expensive, but allows researchers to identify different cells in a population
- Can be very important, for examples, in studying tumours
- Nevertheless, due to amplification processes, etc. there can be biases

Single Cell RNA Sequencing Workflow



Outline

1. Protein Expression
2. **Correlations and Causality**

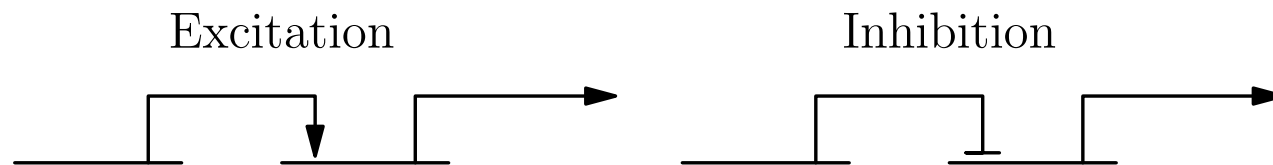


Correlation Between Time Series

- We suppose we have carried out an experiment where we have put a population of cells in a saline solution
- For each mRNA we have a time series (e.g. taken every 30min)

$$\mathbf{X}_i = (X_i(1), X_i(2), \dots, X_i(T))$$

- We can immediately measure if two mRNAs are correlated or anti-correlated

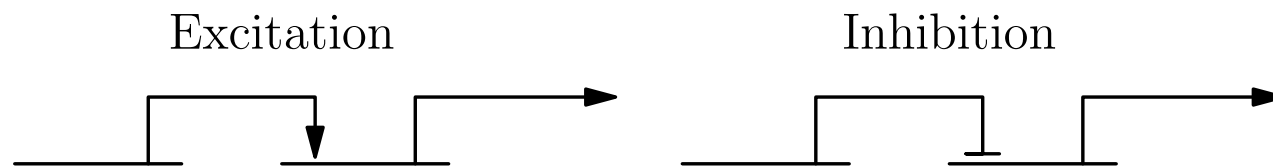


Correlation Between Time Series

- We suppose we have carried out an experiment where we have put a population of cells in a saline solution
- For each mRNA we have a time series (e.g. taken every 30min)

$$\mathbf{X}_i = (X_i(1), X_i(2), \dots, X_i(T))$$

- We can immediately measure if two mRNAs are correlated or anti-correlated

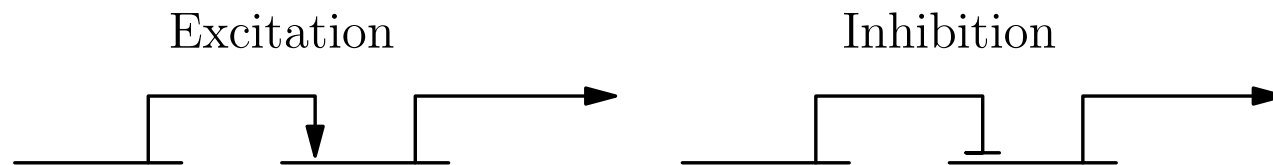


Correlation Between Time Series


- We suppose we have carried out an experiment where we have put a population of cells in a saline solution
- For each mRNA we have a time series (e.g. taken every 30min)

$$\mathbf{X}_i = (X_i(1), X_i(2), \dots, X_i(T))$$

- We can immediately measure if two mRNAs are correlated or anti-correlated



Pearson Correlation



Correlation and Causality

- The famous mantra is

Correlation does not imply causality

- If two random variables X and Y are correlated then
 - ★ X could affect Y
 - ★ Y could affect X
 - ★ X and Y could not influence each other, but both could be affected by another random variable Z

Correlation and Causality

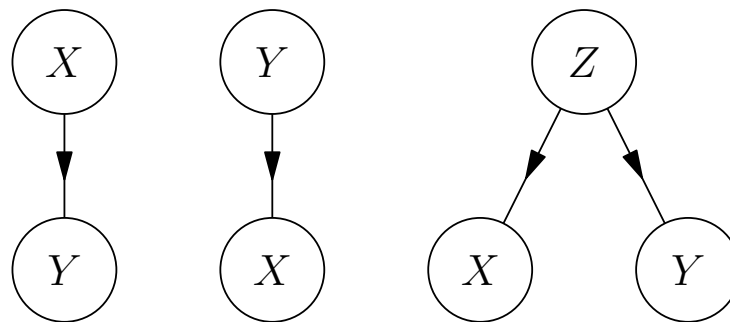
- The famous mantra is

Correlation does not imply causality

- If two random variables X and Y are correlated then
 - ★ X could affect Y
 - ★ Y could affect X
 - ★ X and Y could not influence each other, but both could be affected by another random variable Z

Graphical Models

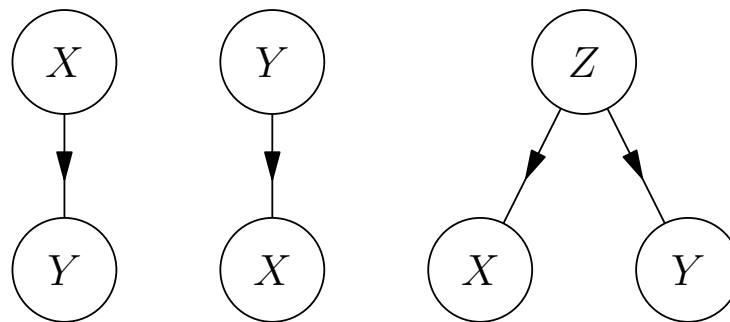
- Graphical models are directed graphs that show causal relationships between random variables
- We could represent the three conditions described above by



- Want to infer a graphical model showing interactions between proteins

Graphical Models

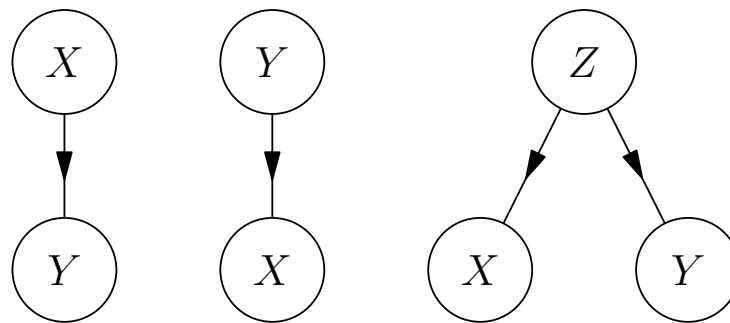
- Graphical models are directed graphs that show causal relationships between random variables
- We could represent the three conditions described above by



- Want to infer a graphical model showing interactions between proteins

Graphical Models

- Graphical models are directed graphs that show causal relationships between random variables
- We could represent the three conditions described above by



- Want to infer a graphical model showing interactions between proteins

Statistical Independence

- Two random variables are statistically independent if

$$\mathbb{P}(X, Y) = \mathbb{P}(X) \mathbb{P}(Y)$$

- Equally this implies $\mathbb{P}(X|Y) = \mathbb{P}(X)$ and $\mathbb{P}(Y|X) = \mathbb{P}(Y)$
- Statistically independent variables are uncorrelated
- But statistical independence is often too powerful

Statistical Independence

- Two random variables are statistically independent if

$$\mathbb{P}(X, Y) = \mathbb{P}(X) \mathbb{P}(Y)$$

- Equally this implies $\mathbb{P}(X|Y) = \mathbb{P}(X)$ and $\mathbb{P}(Y|X) = \mathbb{P}(Y)$
- Statistically independent variables are uncorrelated
- But statistical independence is often too powerful

Statistical Independence

- Two random variables are statistically independent if

$$\mathbb{P}(X, Y) = \mathbb{P}(X) \mathbb{P}(Y)$$

- Equally this implies $\mathbb{P}(X|Y) = \mathbb{P}(X)$ and $\mathbb{P}(Y|X) = \mathbb{P}(Y)$
- Statistically independent variables are uncorrelated
- But statistical independence is often too powerful

Statistical Independence

- Two random variables are statistically independent if


$$\mathbb{P}(X, Y) = \mathbb{P}(X) \mathbb{P}(Y)$$

- Equally this implies $\mathbb{P}(X|Y) = \mathbb{P}(X)$ and $\mathbb{P}(Y|X) = \mathbb{P}(Y)$
- Statistically independent variables are uncorrelated
- But statistical independence is often too powerful

Conditional Independence

- A weaker notion is conditional independence

$$\mathbb{P}(X, Y|Z) = \mathbb{P}(X|Z) \mathbb{P}(Y|Z)$$

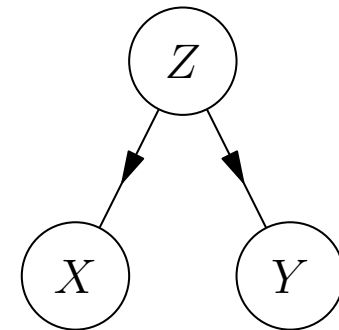


- Conditional independence implies that there is no direct causation
- This is what we want to know
- But it doesn't imply zero correlation
- How do we identify conditionally independent quantities?

Conditional Independence

- A weaker notion is conditional independence

$$\mathbb{P}(X, Y|Z) = \mathbb{P}(X|Z) \mathbb{P}(Y|Z)$$

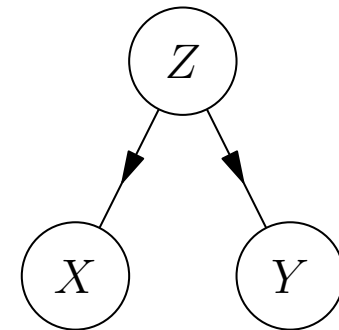


- Conditional independence implies that there is no direct causation
- This is what we want to know
- But it doesn't imply zero correlation
- How do we identify conditionally independent quantities?

Conditional Independence

- A weaker notion is conditional independence

$$\mathbb{P}(X, Y|Z) = \mathbb{P}(X|Z) \mathbb{P}(Y|Z)$$

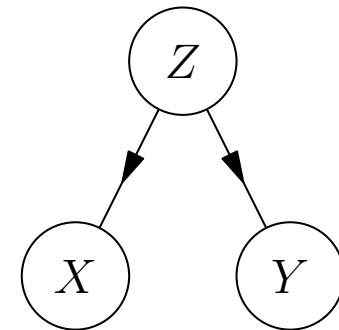


- Conditional independence implies that there is no direct causation
- This is what we want to know
- But it doesn't imply zero correlation
- How do we identify conditionally independent quantities?

Conditional Independence

- A weaker notion is conditional independence

$$\mathbb{P}(X, Y|Z) = \mathbb{P}(X|Z) \mathbb{P}(Y|Z)$$

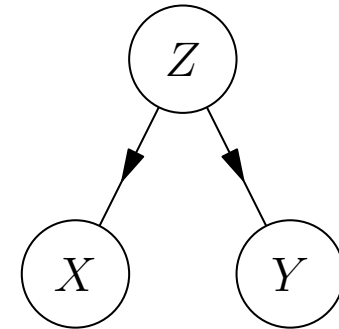


- Conditional independence implies that there is no direct causation
- This is what we want to know
- But it doesn't imply zero correlation
- How do we identify conditionally independent quantities?

Conditional Independence

- A weaker notion is conditional independence

$$\mathbb{P}(X, Y|Z) = \mathbb{P}(X|Z) \mathbb{P}(Y|Z)$$



- Conditional independence implies that there is no direct causation
- This is what we want to know
- But it doesn't imply zero correlation
- How do we identify conditionally independent quantities?

Multivariate Normal Distributions

- The simplest probability distribution for describing continuous variables is the multivariate normal distribution

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} \sqrt{|\boldsymbol{\Sigma}|}} e^{-\frac{1}{2}(\mathbf{X}-\boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{X}-\boldsymbol{\mu})}$$

- With

$$\mathbb{E}[\mathbf{X}] = \int \mathbf{X} \mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) d\mathbf{X} = \boldsymbol{\mu}$$

$$\text{Cov}[\mathbf{X}, \mathbf{X}] = \mathbb{E}[\mathbf{X} \mathbf{X}^\top] - \mathbb{E}[\mathbf{X}] \mathbb{E}[\mathbf{X}^\top] = \boldsymbol{\Sigma}$$

Multivariate Normal Distributions

- The simplest probability distribution for describing continuous variables is the multivariate normal distribution

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} \sqrt{|\boldsymbol{\Sigma}|}} e^{-\frac{1}{2}(\mathbf{X}-\boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{X}-\boldsymbol{\mu})}$$

- With

$$\mathbb{E}[\mathbf{X}] = \int \mathbf{X} \mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) d\mathbf{X} = \boldsymbol{\mu}$$

$$\text{Cov}[\mathbf{X}, \mathbf{X}] = \mathbb{E}[\mathbf{X} \mathbf{X}^\top] - \mathbb{E}[\mathbf{X}] \mathbb{E}[\mathbf{X}^\top] = \boldsymbol{\Sigma}$$

Multivariate Normal Distributions

- The simplest probability distribution for describing continuous variables is the multivariate normal distribution

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} \sqrt{|\boldsymbol{\Sigma}|}} e^{-\frac{1}{2}(\mathbf{X}-\boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{X}-\boldsymbol{\mu})}$$

- With

$$\mathbb{E}[\mathbf{X}] = \int \mathbf{X} \mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) d\mathbf{X} = \boldsymbol{\mu}$$

$$\text{Cov}[\mathbf{X}, \mathbf{X}] = \mathbb{E}[\mathbf{X} \mathbf{X}^\top] - \mathbb{E}[\mathbf{X}] \mathbb{E}[\mathbf{X}^\top] = \boldsymbol{\Sigma}$$

Covariance

- The covariance $\mathbb{Cov}[X_i, X_j]$ measures how X_i varies with X_j
- The correlation is given by

$$q_{ij} = \frac{\mathbb{Cov}[X_i, X_j]}{\sqrt{\mathbb{Var}[X_i] \mathbb{Var}[X_j]}}$$

- Where $\mathbb{Var}[X_i] = \mathbb{Cov}[X_i, X_i] = \mathbb{E}[X_i^2] - \mathbb{E}[X_i]^2$ is the variance of X_i

Covariance

- The covariance $\mathbb{Cov}[X_i, X_j]$ measures how X_i varies with X_j
- The correlation is given by

$$q_{ij} = \frac{\mathbb{Cov}[X_i, X_j]}{\sqrt{\mathbb{Var}[X_i] \mathbb{Var}[X_j]}}$$

- Where $\mathbb{Var}[X_i] = \mathbb{Cov}[X_i, X_i] = \mathbb{E}[X_i^2] - \mathbb{E}[X_i]^2$ is the variance of X_i

Trivariate Normal Distribution

- Consider a 3-D normally distributed random variable $\mathbf{X} = (X_1, X_2, X_3)$

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|} \exp \left(-\frac{1}{2} \Sigma_{11}^{-1} (X_1 - \mu_1)^2 - \frac{1}{2} \Sigma_{22}^{-1} (X_2 - \mu_2)^2 \right. \\ \left. - \frac{1}{2} \Sigma_{33}^{-1} (X_3 - \mu_3)^2 - \Sigma_{12}^{-1} (X_1 - \mu_1) (X_2 - \mu_2) \right. \\ \left. - \Sigma_{23}^{-1} (X_2 - \mu_2) (X_3 - \mu_3) - \Sigma_{13}^{-1} (X_1 - \mu_1) (X_3 - \mu_3) \right)$$

- If $\Sigma_{13}^{-1} = 0$ then

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) \propto e^{-\frac{1}{2} \Sigma_{11}^{-1} (X_1 - \mu_1)^2 - \Sigma_{12}^{-1} (X_1 - \mu_1) (X_2 - \mu_2)} \\ e^{-\frac{1}{2} \Sigma_{33}^{-1} (X_3 - \mu_3)^2 - \Sigma_{23}^{-1} (X_2 - \mu_2) (X_3 - \mu_3)} e^{-\frac{1}{2} \Sigma_{22}^{-1} (X_2 - \mu_2)^2}$$

Trivariate Normal Distribution

- Consider a 3-D normally distributed random variable $\mathbf{X} = (X_1, X_2, X_3)$

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|} \exp \left(-\frac{1}{2} \Sigma_{11}^{-1} (X_1 - \mu_1)^2 - \frac{1}{2} \Sigma_{22}^{-1} (X_2 - \mu_2)^2 \right. \\ \left. - \frac{1}{2} \Sigma_{33}^{-1} (X_3 - \mu_3)^2 - \Sigma_{12}^{-1} (X_1 - \mu_1) (X_2 - \mu_2) \right. \\ \left. - \Sigma_{23}^{-1} (X_2 - \mu_2) (X_3 - \mu_3) - \Sigma_{13}^{-1} (X_1 - \mu_1) (X_3 - \mu_3) \right)$$

- If $\Sigma_{13}^{-1} = 0$ then

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) \propto e^{-\frac{1}{2} \Sigma_{11}^{-1} (X_1 - \mu_1)^2 - \Sigma_{12}^{-1} (X_1 - \mu_1) (X_2 - \mu_2)} \\ e^{-\frac{1}{2} \Sigma_{33}^{-1} (X_3 - \mu_3)^2 - \Sigma_{23}^{-1} (X_2 - \mu_2) (X_3 - \mu_3)} e^{-\frac{1}{2} \Sigma_{22}^{-1} (X_2 - \mu_2)^2}$$

Conditional Independence

- If $\Sigma_{13}^{-1} = 0$ then

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) \propto e^{-\frac{1}{2} \Sigma_{11}^{-1} (X_1 - \mu_1)^2 - \Sigma_{12}^{-1} (X_1 - \mu_1) (X_2 - \mu_2)} \\ e^{-\frac{1}{2} \Sigma_{33}^{-1} (X_3 - \mu_3)^2 - \Sigma_{23}^{-1} (X_2 - \mu_2) (X_3 - \mu_3)} e^{-\frac{1}{2} \Sigma_{22}^{-1} (X_2 - \mu_2)^2}$$

- But $\mathbb{P}(X_1, X_3|X_2) = \mathbb{P}(X_1, X_2, X_3) / \mathbb{P}(X_2)$
- So

$$\mathbb{P}(X_1, X_3|X_2) = \mathbb{P}(X_1|X_2) \mathbb{P}(X_3|X_2)$$

- Thus $\Sigma_{13}^{-1} = 0 \Rightarrow X_1 \perp X_3|X_2$

Conditional Independence

- If $\Sigma_{13}^{-1} = 0$ then

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) \propto e^{-\frac{1}{2} \Sigma_{11}^{-1} (X_1 - \mu_1)^2 - \Sigma_{12}^{-1} (X_1 - \mu_1) (X_2 - \mu_2)} \\ e^{-\frac{1}{2} \Sigma_{33}^{-1} (X_3 - \mu_3)^2 - \Sigma_{23}^{-1} (X_2 - \mu_2) (X_3 - \mu_3)} e^{-\frac{1}{2} \Sigma_{22}^{-1} (X_2 - \mu_2)^2}$$

- But $\mathbb{P}(X_1, X_3|X_2) = \mathbb{P}(X_1, X_2, X_3) / \mathbb{P}(X_2)$

- So

$$\mathbb{P}(X_1, X_3|X_2) = \mathbb{P}(X_1|X_2) \mathbb{P}(X_3|X_2)$$

- Thus $\Sigma_{13}^{-1} = 0 \Rightarrow X_1 \perp X_3|X_2$

Conditional Independence

- If $\Sigma_{13}^{-1} = 0$ then

$$\mathcal{N}(\mathbf{X}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) \propto e^{-\frac{1}{2} \Sigma_{11}^{-1} (X_1 - \mu_1)^2 - \Sigma_{12}^{-1} (X_1 - \mu_1) (X_2 - \mu_2)} \\ e^{-\frac{1}{2} \Sigma_{33}^{-1} (X_3 - \mu_3)^2 - \Sigma_{23}^{-1} (X_2 - \mu_2) (X_3 - \mu_3)} e^{-\frac{1}{2} \Sigma_{22}^{-1} (X_2 - \mu_2)^2}$$

- But $\mathbb{P}(X_1, X_3|X_2) = \mathbb{P}(X_1, X_2, X_3) / \mathbb{P}(X_2)$
- So

$$\mathbb{P}(X_1, X_3|X_2) = \mathbb{P}(X_1|X_2) \mathbb{P}(X_3|X_2)$$

- Thus $\Sigma_{13}^{-1} = 0 \Rightarrow X_1 \perp X_3|X_2$

Finding Conditional Independence

- If we suppose that the expression data at different times is normally distributed
- One way to find conditionally independent variables would be to measure the covariance between the RNA levels
 $\text{Cov}[X_i, X_j] = \Sigma_{ij}$ and invert the matrix Σ
- But we can only estimate the empirical covariance

$$\hat{\Sigma}_{ij} = \frac{1}{T-1} \sum_{t=1}^T (X_i(t) - \hat{\mu}_i) (X_j(t) - \hat{\mu}_j), \quad \hat{\mu}_i = \frac{1}{T} \sum_{t=1}^T X_i(t)$$

- How do we tell if Σ_{ij}^{-1} is zero or just small?

Finding Conditional Independence

- If we suppose that the expression data at different times is normally distributed—**which of course it isn't**
- One way to find conditionally independent variables would be to measure the covariance between the RNA levels
 $\text{Cov}[X_i, X_j] = \Sigma_{ij}$ and invert the matrix Σ
- But we can only estimate the empirical covariance

$$\hat{\Sigma}_{ij} = \frac{1}{T-1} \sum_{t=1}^T (X_i(t) - \hat{\mu}_i) (X_j(t) - \hat{\mu}_j), \quad \hat{\mu}_i = \frac{1}{T} \sum_{t=1}^T X_i(t)$$

- How do we tell if Σ_{ij}^{-1} is zero or just small?

Finding Conditional Independence

- If we suppose that the expression data at different times is normally distributed—which of course it isn't **quite**
- One way to find conditionally independent variables would be to measure the covariance between the RNA levels
 $\text{Cov}[X_i, X_j] = \Sigma_{ij}$ and invert the matrix Σ
- But we can only estimate the empirical covariance

$$\hat{\Sigma}_{ij} = \frac{1}{T-1} \sum_{t=1}^T (X_i(t) - \hat{\mu}_i) (X_j(t) - \hat{\mu}_j), \quad \hat{\mu}_i = \frac{1}{T} \sum_{t=1}^T X_i(t)$$

- How do we tell if Σ_{ij}^{-1} is zero or just small?

Finding Conditional Independence

- If we suppose that the expression data at different times is normally distributed—which of course it isn't quite
- One way to find conditionally independent variables would be to measure the covariance between the RNA levels
 $\text{Cov}[X_i, X_j] = \Sigma_{ij}$ and invert the matrix Σ
- But we can only estimate the empirical covariance

$$\hat{\Sigma}_{ij} = \frac{1}{T-1} \sum_{t=1}^T (X_i(t) - \hat{\mu}_i) (X_j(t) - \hat{\mu}_j), \quad \hat{\mu}_i = \frac{1}{T} \sum_{t=1}^T X_i(t)$$

- How do we tell if Σ_{ij}^{-1} is zero or just small?

Finding Conditional Independence

- If we suppose that the expression data at different times is normally distributed—which of course it isn't quite
- One way to find conditionally independent variables would be to measure the covariance between the RNA levels $\text{Cov}[X_i, X_j] = \Sigma_{ij}$ and invert the matrix Σ
- But we can only estimate the empirical covariance

$$\hat{\Sigma}_{ij} = \frac{1}{T-1} \sum_{t=1}^T (X_i(t) - \hat{\mu}_i) (X_j(t) - \hat{\mu}_j), \quad \hat{\mu}_i = \frac{1}{T} \sum_{t=1}^T X_i(t)$$

- How do we tell if Σ_{ij}^{-1} is zero or just small?

Finding Conditional Independence

- If we suppose that the expression data at different times is normally distributed—which of course it isn't quite
- One way to find conditionally independent variables would be to measure the covariance between the RNA levels $\text{Cov}[X_i, X_j] = \Sigma_{ij}$ and invert the matrix Σ
- But we can only estimate the empirical covariance

$$\hat{\Sigma}_{ij} = \frac{1}{T-1} \sum_{t=1}^T (X_i(t) - \hat{\mu}_i) (X_j(t) - \hat{\mu}_j), \quad \hat{\mu}_i = \frac{1}{T} \sum_{t=1}^T X_i(t)$$

- How do we tell if Σ_{ij}^{-1} is zero or just small?

Maximum Likelihood

- Given our data $\mathcal{D} = (X_i(t) | i = 1, 2, \dots, n \quad t = 1, 2, \dots, T)$ we can try to find the value of $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ that maximises the likelihood of the data

$$\begin{aligned}\mathbb{P}(\mathcal{D} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) &= \prod_{t=1}^T \mathcal{N}(\mathbf{X}(t) | \boldsymbol{\mu}, \boldsymbol{\Sigma}) \\ &= \left(\frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \right)^T \prod_{t=1}^T e^{-\frac{1}{2} (\mathbf{X}(t) - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{X}(t) - \boldsymbol{\mu})}\end{aligned}$$

- Where $\mathbf{X}(t) = (X_1(t), X_2(t), \dots, X_n(t))$ (note that when looking at correlation we considered $\mathbf{X}_i$ as vectors)

Maximum Likelihood

- Given our data $\mathcal{D} = (X_i(t) | i = 1, 2, \dots, n \quad t = 1, 2, \dots, T)$ we can try to find the value of $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ that maximises the likelihood of the data

$$\begin{aligned}\mathbb{P}(\mathcal{D} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) &= \prod_{t=1}^T \mathcal{N}(\mathbf{X}(t) | \boldsymbol{\mu}, \boldsymbol{\Sigma}) \\ &= \left(\frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \right)^T \prod_{t=1}^T e^{-\frac{1}{2} (\mathbf{X}(t) - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{X}(t) - \boldsymbol{\mu})}\end{aligned}$$

- Where $\mathbf{X}(t) = (X_1(t), X_2(t), \dots, X_n(t))$ (note that when looking at correlation we considered $\mathbf{X}_i$ as vectors)

Maximum Likelihood

- Given our data $\mathcal{D} = (X_i(t) | i = 1, 2, \dots, n \quad t = 1, 2, \dots, T)$ we can try to find the value of $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ that maximises the likelihood of the data

$$\begin{aligned}\mathbb{P}(\mathcal{D} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) &= \prod_{t=1}^T \mathcal{N}(\mathbf{X}(t) | \boldsymbol{\mu}, \boldsymbol{\Sigma}) \\ &= \left(\frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \right)^T \prod_{t=1}^T e^{-\frac{1}{2} (\mathbf{X}(t) - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{X}(t) - \boldsymbol{\mu})}\end{aligned}$$

- Where $\mathbf{X}(t) = (X_1(t), X_2(t), \dots, X_n(t))$ (note that when looking at correlation we considered $\mathbf{X}_i$ as vectors)

Maximum Log-Likelihood

- The maximum of the likelihood will also be the maximum of the log-likelihood $\mathcal{L} = \log(\mathbb{P}(\mathcal{D}|\boldsymbol{\mu}, \boldsymbol{\Sigma}))$

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \sum_{t=1}^T (\mathbf{X}(t) - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{X}(t) - \boldsymbol{\mu}) + C$$

- If we maximise with respect to $\boldsymbol{\mu}$ (i.e. solving $\nabla_{\boldsymbol{\mu}} \mathcal{L} = 0$) we find $\boldsymbol{\mu}^* = \hat{\boldsymbol{\mu}}$
- Using the maximum likelihood mean

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \text{tr} \mathbf{S} \boldsymbol{\Sigma}^{-1}$$

- with $\mathbf{S} = \frac{1}{T} \sum_{t=1}^T (\mathbf{X}(t) - \hat{\boldsymbol{\mu}}) (\mathbf{X}(t) - \hat{\boldsymbol{\mu}})^\top$

Maximum Log-Likelihood

- The maximum of the likelihood will also be the maximum of the log-likelihood $\mathcal{L} = \log(\mathbb{P}(\mathcal{D}|\boldsymbol{\mu}, \boldsymbol{\Sigma}))$

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \sum_{t=1}^T (\mathbf{X}(t) - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{X}(t) - \boldsymbol{\mu}) + C$$

- If we maximise with respect to $\boldsymbol{\mu}$ (i.e. solving $\nabla_{\boldsymbol{\mu}} \mathcal{L} = 0$) we find $\boldsymbol{\mu}^* = \hat{\boldsymbol{\mu}}$

- Using the maximum likelihood mean

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \text{tr} \mathbf{S} \boldsymbol{\Sigma}^{-1}$$

- with $\mathbf{S} = \frac{1}{T} \sum_{t=1}^T (\mathbf{X}(t) - \hat{\boldsymbol{\mu}}) (\mathbf{X}(t) - \hat{\boldsymbol{\mu}})^\top$

Maximum Log-Likelihood

- The maximum of the likelihood will also be the maximum of the log-likelihood $\mathcal{L} = \log(\mathbb{P}(\mathcal{D}|\boldsymbol{\mu}, \boldsymbol{\Sigma}))$

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \sum_{t=1}^T (\mathbf{X}(t) - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{X}(t) - \boldsymbol{\mu}) + C$$

- If we maximise with respect to $\boldsymbol{\mu}$ (i.e. solving $\nabla_{\boldsymbol{\mu}} \mathcal{L} = 0$) we find $\boldsymbol{\mu}^* = \hat{\boldsymbol{\mu}}$

- Using the maximum likelihood mean

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \text{tr} \mathbf{S} \boldsymbol{\Sigma}^{-1}$$

- with $\mathbf{S} = \frac{1}{T} \sum_{t=1}^T (\mathbf{X}(t) - \hat{\boldsymbol{\mu}}) (\mathbf{X}(t) - \hat{\boldsymbol{\mu}})^\top$

Maximum Log-Likelihood

- The maximum of the likelihood will also be the maximum of the log-likelihood $\mathcal{L} = \log(\mathbb{P}(\mathcal{D}|\boldsymbol{\mu}, \boldsymbol{\Sigma}))$

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \sum_{t=1}^T (\mathbf{X}(t) - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{X}(t) - \boldsymbol{\mu}) + C$$

- If we maximise with respect to $\boldsymbol{\mu}$ (i.e. solving $\nabla_{\boldsymbol{\mu}} \mathcal{L} = 0$) we find $\boldsymbol{\mu}^* = \hat{\boldsymbol{\mu}}$
- Using the maximum likelihood mean

$$\mathcal{L} = -\frac{T}{2} \log(|\boldsymbol{\Sigma}|) - \frac{1}{2} \text{tr} \mathbf{S} \boldsymbol{\Sigma}^{-1}$$

- with $\mathbf{S} = \frac{1}{T} \sum_{t=1}^T (\mathbf{X}(t) - \hat{\boldsymbol{\mu}}) (\mathbf{X}(t) - \hat{\boldsymbol{\mu}})^\top$

Regularisation

- The maximum likelihood value of the covariance is (solving for $\partial\mathcal{L}/\partial\mathbf{\Sigma} = 0$) given by $\mathbf{\Sigma} = \mathbf{S}$
- However we want to find solutions where $\mathbf{\Theta} = \mathbf{\Sigma}^{-1}$ is sparse
- Note that $|\mathbf{\Sigma}| = 1/|\mathbf{\Sigma}^{-1}| = 1/|\mathbf{\Theta}|$
- So $\mathcal{L} = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta}$
- We can add a “regulariser” to drive small Θ_{ij} to 0

$$\mathcal{L} - \|\mathbf{\Theta}\|_1 = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta} - \|\mathbf{\Theta}\|_1$$

Regularisation

- The maximum likelihood value of the covariance is (solving for $\partial\mathcal{L}/\partial\mathbf{\Sigma} = 0$) given by $\mathbf{\Sigma} = \mathbf{S}$
- However we want to find solutions where $\mathbf{\Theta} = \mathbf{\Sigma}^{-1}$ is sparse
- Note that $|\mathbf{\Sigma}| = 1/|\mathbf{\Sigma}^{-1}| = 1/|\mathbf{\Theta}|$
- So $\mathcal{L} = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta}$
- We can add a “regulariser” to drive small Θ_{ij} to 0

$$\mathcal{L} - \|\mathbf{\Theta}\|_1 = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta} - \|\mathbf{\Theta}\|_1$$

Regularisation

- The maximum likelihood value of the covariance is (solving for $\partial\mathcal{L}/\partial\mathbf{\Sigma} = 0$) given by $\mathbf{\Sigma} = \mathbf{S}$
- However we want to find solutions where $\mathbf{\Theta} = \mathbf{\Sigma}^{-1}$ is sparse
- Note that $|\mathbf{\Sigma}| = 1/|\mathbf{\Sigma}^{-1}| = 1/|\mathbf{\Theta}|$
- So $\mathcal{L} = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta}$
- We can add a “regulariser” to drive small Θ_{ij} to 0

$$\mathcal{L} - \|\mathbf{\Theta}\|_1 = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta} - \|\mathbf{\Theta}\|_1$$

Regularisation

- The maximum likelihood value of the covariance is (solving for $\partial\mathcal{L}/\partial\mathbf{\Sigma} = 0$) given by $\mathbf{\Sigma} = \mathbf{S}$
- However we want to find solutions where $\mathbf{\Theta} = \mathbf{\Sigma}^{-1}$ is sparse
- Note that $|\mathbf{\Sigma}| = 1/|\mathbf{\Sigma}^{-1}| = 1/|\mathbf{\Theta}|$
- So $\mathcal{L} = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta}$
- We can add a “regulariser” to drive small Θ_{ij} to 0

$$\mathcal{L} - \|\mathbf{\Theta}\|_1 = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta} - \|\mathbf{\Theta}\|_1$$

Regularisation

- The maximum likelihood value of the covariance is (solving for $\partial\mathcal{L}/\partial\mathbf{\Sigma} = 0$) given by $\mathbf{\Sigma} = \mathbf{S}$
- However we want to find solutions where $\mathbf{\Theta} = \mathbf{\Sigma}^{-1}$ is sparse
- Note that $|\mathbf{\Sigma}| = 1/|\mathbf{\Sigma}^{-1}| = 1/|\mathbf{\Theta}|$
- So $\mathcal{L} = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta}$
- We can add a “regulariser” to drive small Θ_{ij} to 0

$$\mathcal{L} - \|\mathbf{\Theta}\|_1 = \frac{T}{2} \log(|\mathbf{\Theta}|) - \frac{1}{2} \text{tr} \mathbf{S} \mathbf{\Theta} - \|\mathbf{\Theta}\|_1$$

L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2$$

- It encourages finding sparse solutions

L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions

L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$


- It encourages finding sparse solutions

L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser

- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$

- It encourages finding sparse solutions




L_1 Regulariser


- L_1 regulariser are used in regression where it is known as Lasso

$$L = \sum_{k=1}^m (\mathbf{w}^\top \mathbf{x}_k - y_k)^2 + \nu \sum_{i=1}^p |w_i|$$


- It encourages finding sparse solutions



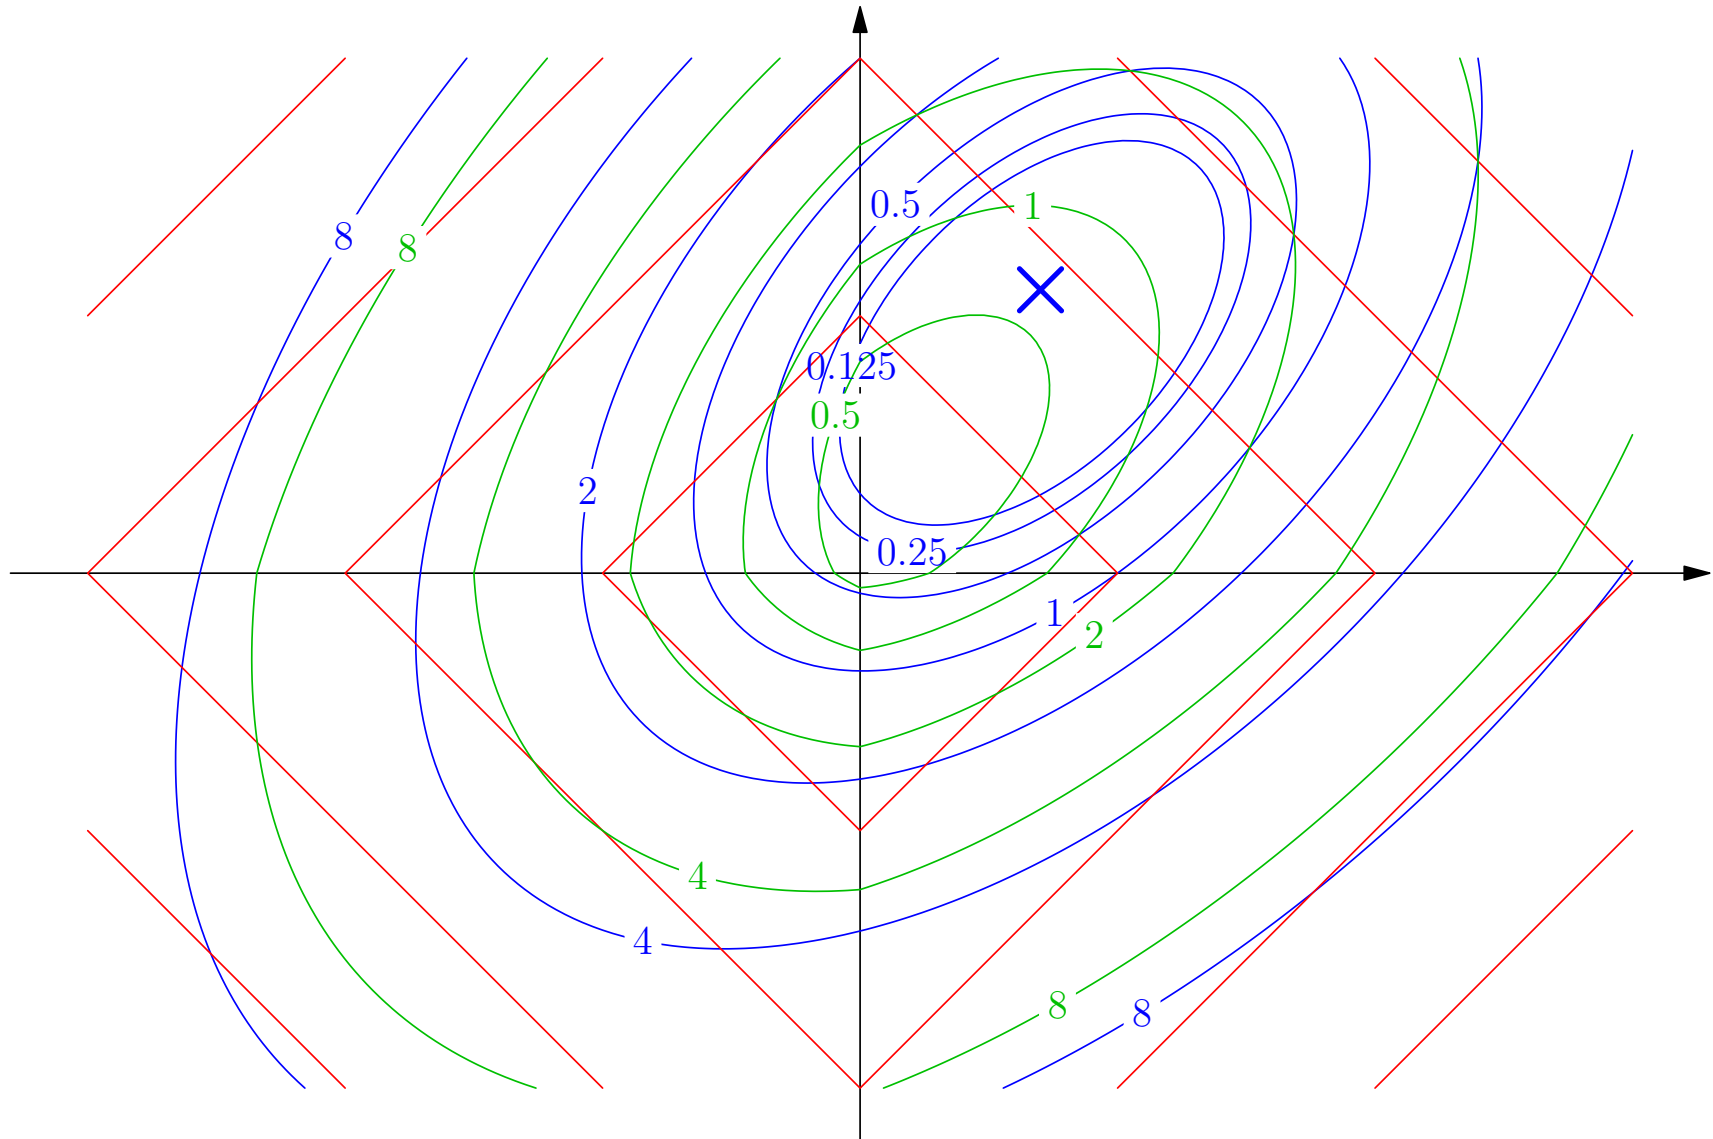
Graphical Lasso



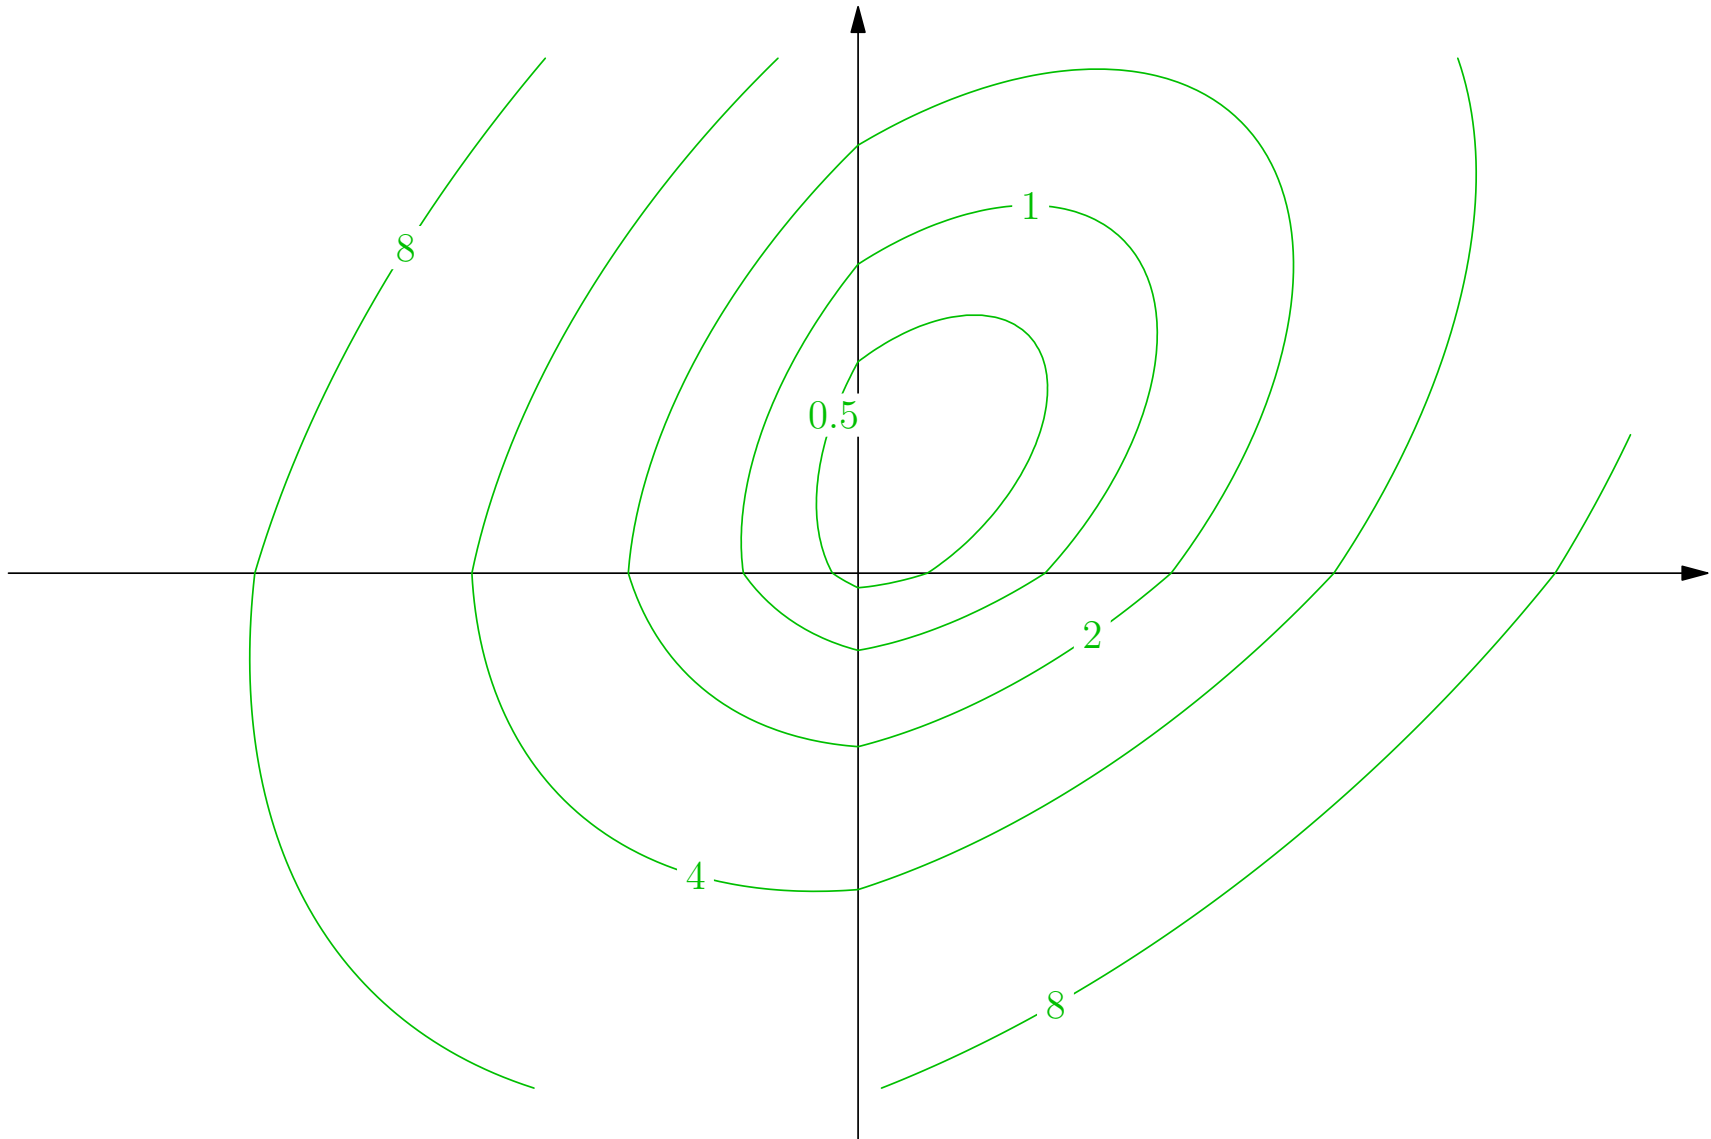
Graphical Lasso




Graphical Lasso




Graphical Lasso



Graphical Lasso



Gene Regulation Network



Summary

- We can use DNA microarrays or RNA sequencing to measure the expression of different proteins
- By measuring the expression over time we can measure correlations between proteins
- To infer a sparse graphical model that explains the observed correlations we can use the “graphical Lasso” to find a sparse inverse correlation matrix
- This is a first step to understand the control mechanisms of a cell

Summary

- We can use DNA microarrays or RNA sequencing to measure the expression of different proteins
- By measuring the expression over time we can measure correlations between proteins
- To infer a sparse graphical model that explains the observed correlations we can use the “graphical Lasso” to find a sparse inverse correlation matrix
- This is a first step to understand the control mechanisms of a cell

Summary

- We can use DNA microarrays or RNA sequencing to measure the expression of different proteins
- By measuring the expression over time we can measure correlations between proteins
- To infer a sparse graphical model that explains the observed correlations we can use the “graphical Lasso” to find a sparse inverse correlation matrix
- This is a first step to understand the control mechanisms of a cell

Summary

- We can use DNA microarrays or RNA sequencing to measure the expression of different proteins
- By measuring the expression over time we can measure correlations between proteins
- To infer a sparse graphical model that explains the observed correlations we can use the “graphical Lasso” to find a sparse inverse correlation matrix
- This is a first step to understand the control mechanisms of a cell